

Task Category	Frontier (Closed)	Fast / Cheap Closed	Open Weights
Complex Reasoning and Math	Top pick. Opus 4.7 (HLE, GPQA leader); GPT-5.5 close. Gemini 3.1 Pro strong + cheaper.	Workable for moderate problems. Haiku 4.5 / GPT-5 mini handle most school-level tasks.	Kimi K2.6 (HLE 54% w/ tools, AIME 96.4) is competitive with frontier here.
Coding and Software Engineering	Opus 4.7 leads SWE-Bench Pro (64.3%); GPT-5.5 leads Terminal-Bench. Pick by workflow.	GPT-5.4 mini & Haiku 4.5 fine for snippets, refactors, sub-agents.	Kimi K2.6 is the only Tier-A open coder; DeepSeek V4 strong at low cost.
Agentic / Long-Horizon Tool Use	GPT-5.5 leads (Terminal-Bench 82.7%, OSWorld). Opus 4.7 leads MCP-Atlas / instruction adherence.	Haiku 4.5: zero tool-call failures in tests; great sub-agent.	Kimi K2.6: 4,000+ tool calls / 12-hr runs; agent swarms.
Long-Form Writing and Drafting	Opus 4.7 (verbose, structured); GPT-5.5 (concise); Gemini 3.1 Pro (longest context).	Haiku 4.5 / Gemini 3 Flash for blog posts, emails, drafts at scale.	Kimi K2.6, Qwen 3.6, Llama-class models — quality near frontier for prose.
Summarization and Documentation Q&A	Gemini 3.1 Pro wins on long docs (1M–2M context). Opus 4.7 / GPT-5.5 also 1M.	Gemini 3.1 Flash-Lite (\$0.25/\$1.50) — best \$ per page processed.	DeepSeek V4 cheapest at scale; Kimi K2.6 for 256K-context multimodal.
Research and Real-Time Web Search	Use frontier models with web tools enabled (all support it now).	Avoid — small models hallucinate more on freshness.	Generally lacks built-in retrieval; pair with a search wrapper.
Image Generation	GPT-5.5 (omnimodal) can call GPT Image 2 inline; Gemini calls Nano Banana 2.	N/A	Flux 2 Dev / Schnell — open weights, self-hostable, near-Pro quality.
Multimodal (Vision / Motion / Audio)	GPT-5.5 (natively omnimodal, single system); Gemini 3.1 Pro (video + 1M ctx).	Gemini 3 Flash for cheap vision tasks at scale.	Kimi K2.6 supports image + video input; Qwen-VL family for OCR/charts.

High Volume / Cost-Sensitive Tasks	Avoid for bulk — too expensive (\$5/\$25+ per 1M).	Top pick. GPT-5.4 mini (\$0.75/\$4.50), Haiku 4.5 (\$1/\$5 with 90% cache discount), Gemini 3.1 Flash-Lite (\$0.25/\$1.50).	DeepSeek V4 (\$0.14/\$0.50) — cheapest viable model; Kimi K2.6 self-hosted.
Private / Local / On-Prem Use	Not available — all closed APIs.	Not available — all closed APIs.	Top pick. Kimi K2.6 (1T MoE, 32B active), DeepSeek V4, Qwen 3.6, Llama, GLM 5.1, Mistral. All have downloadable weights.